



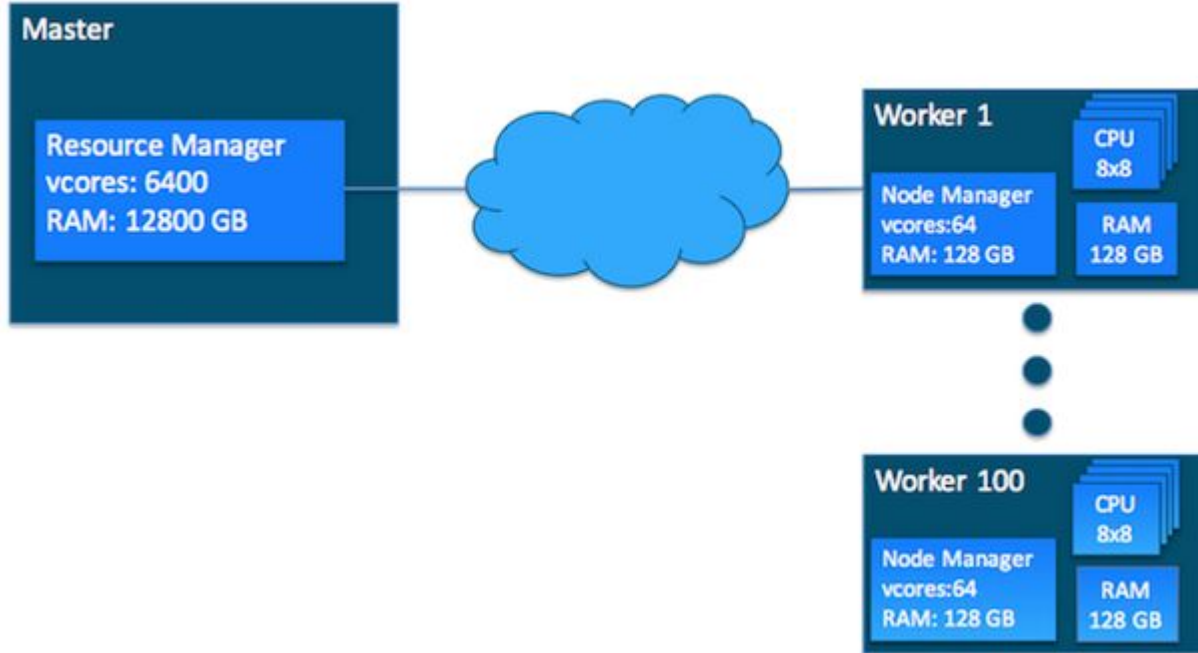
General Cluster Architecture



Resource Manager and Node manager




Vcores and Memory



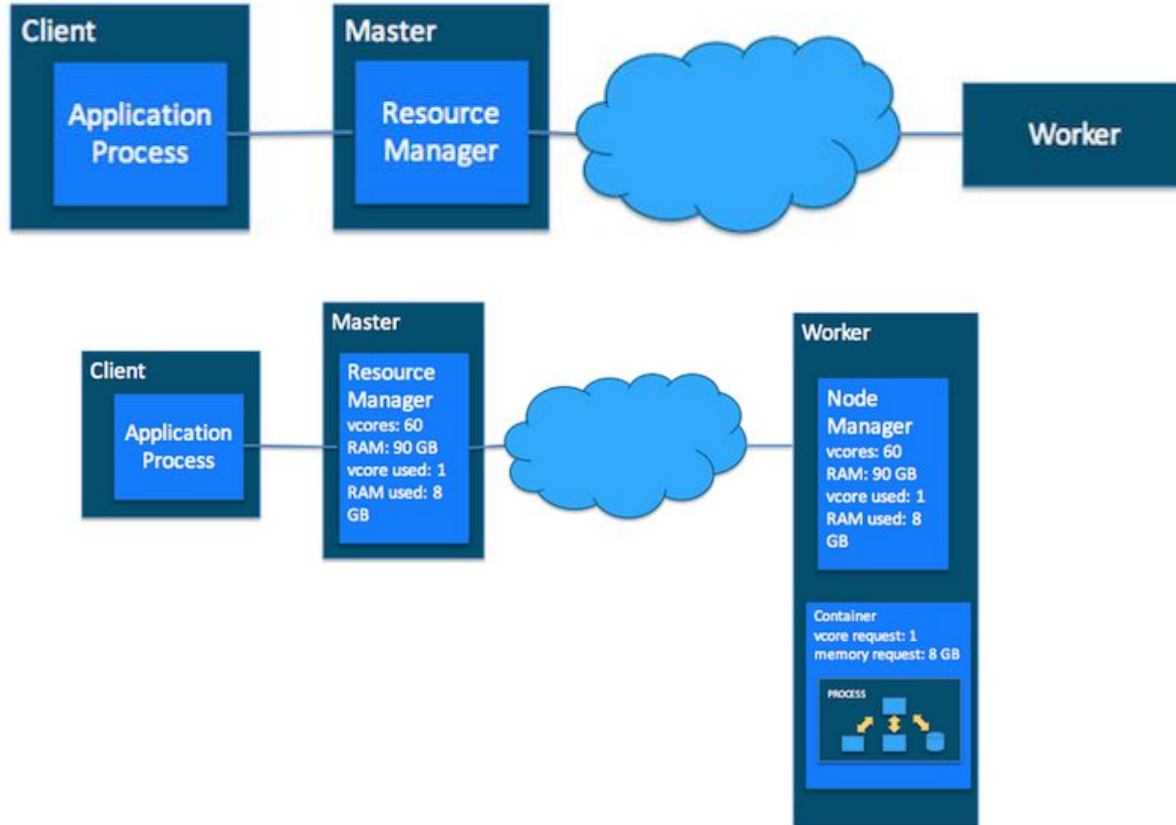
Container, Request, Task, Process

Container
vcore request: 1
memory request: 8 GB

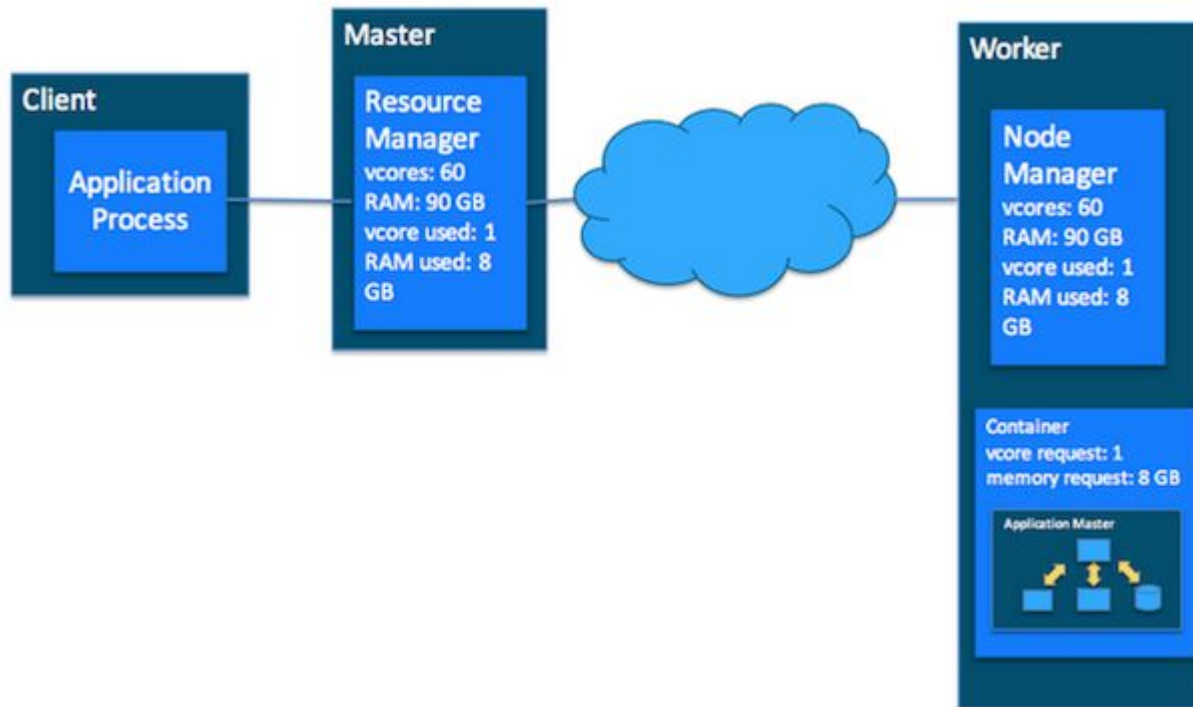
Container
vcore request: 1
memory request: 8 GB



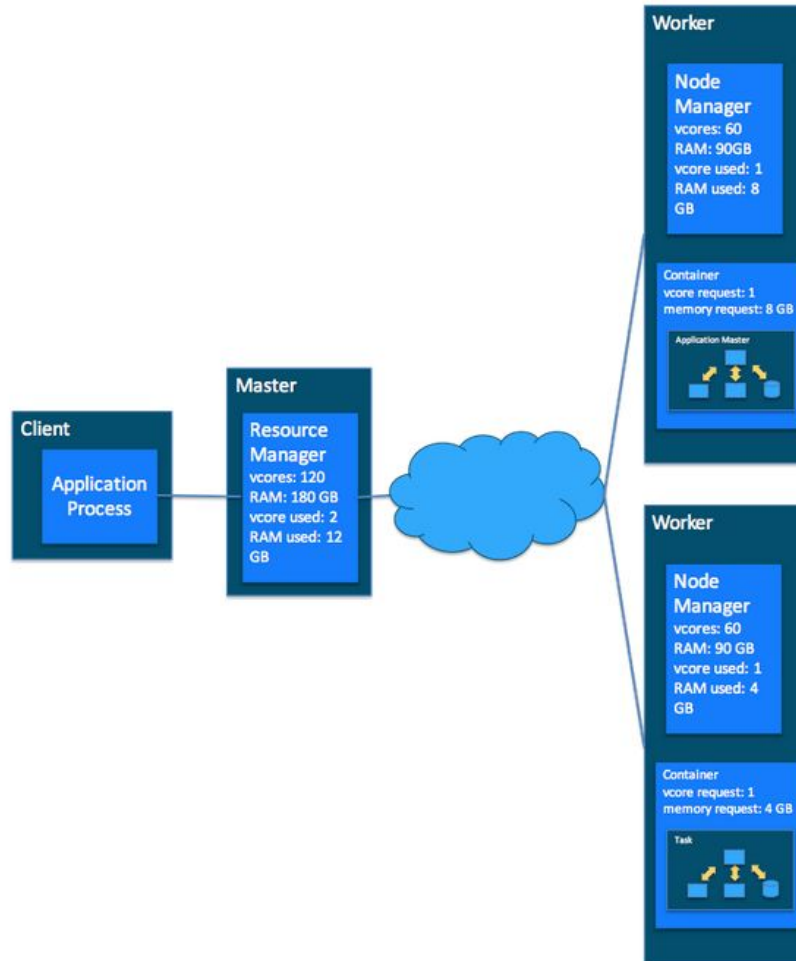
Application Starting



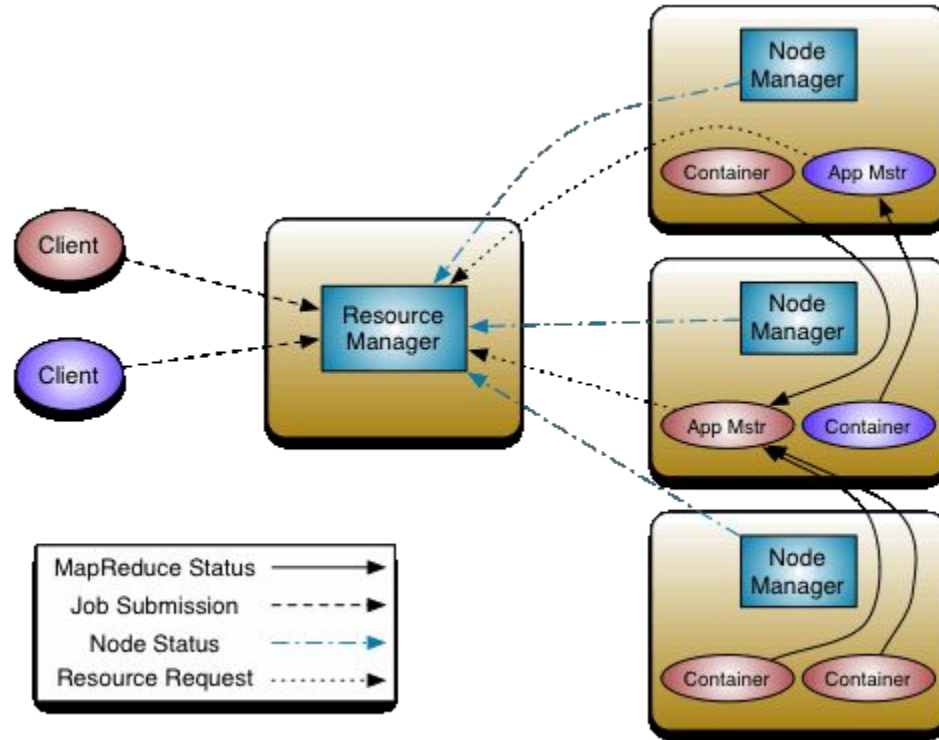
Application Master Container



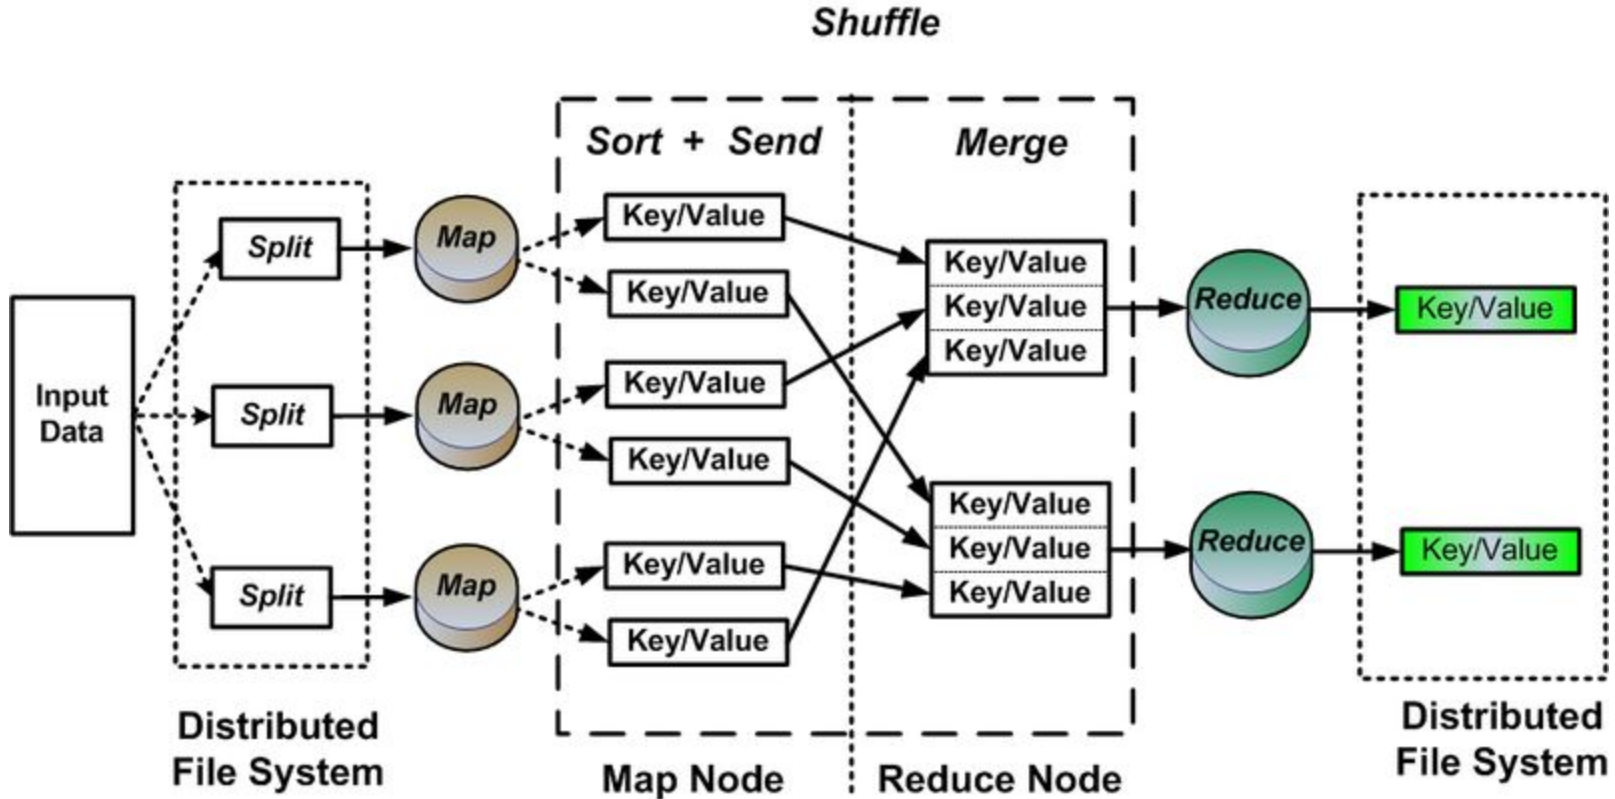
Task Containers

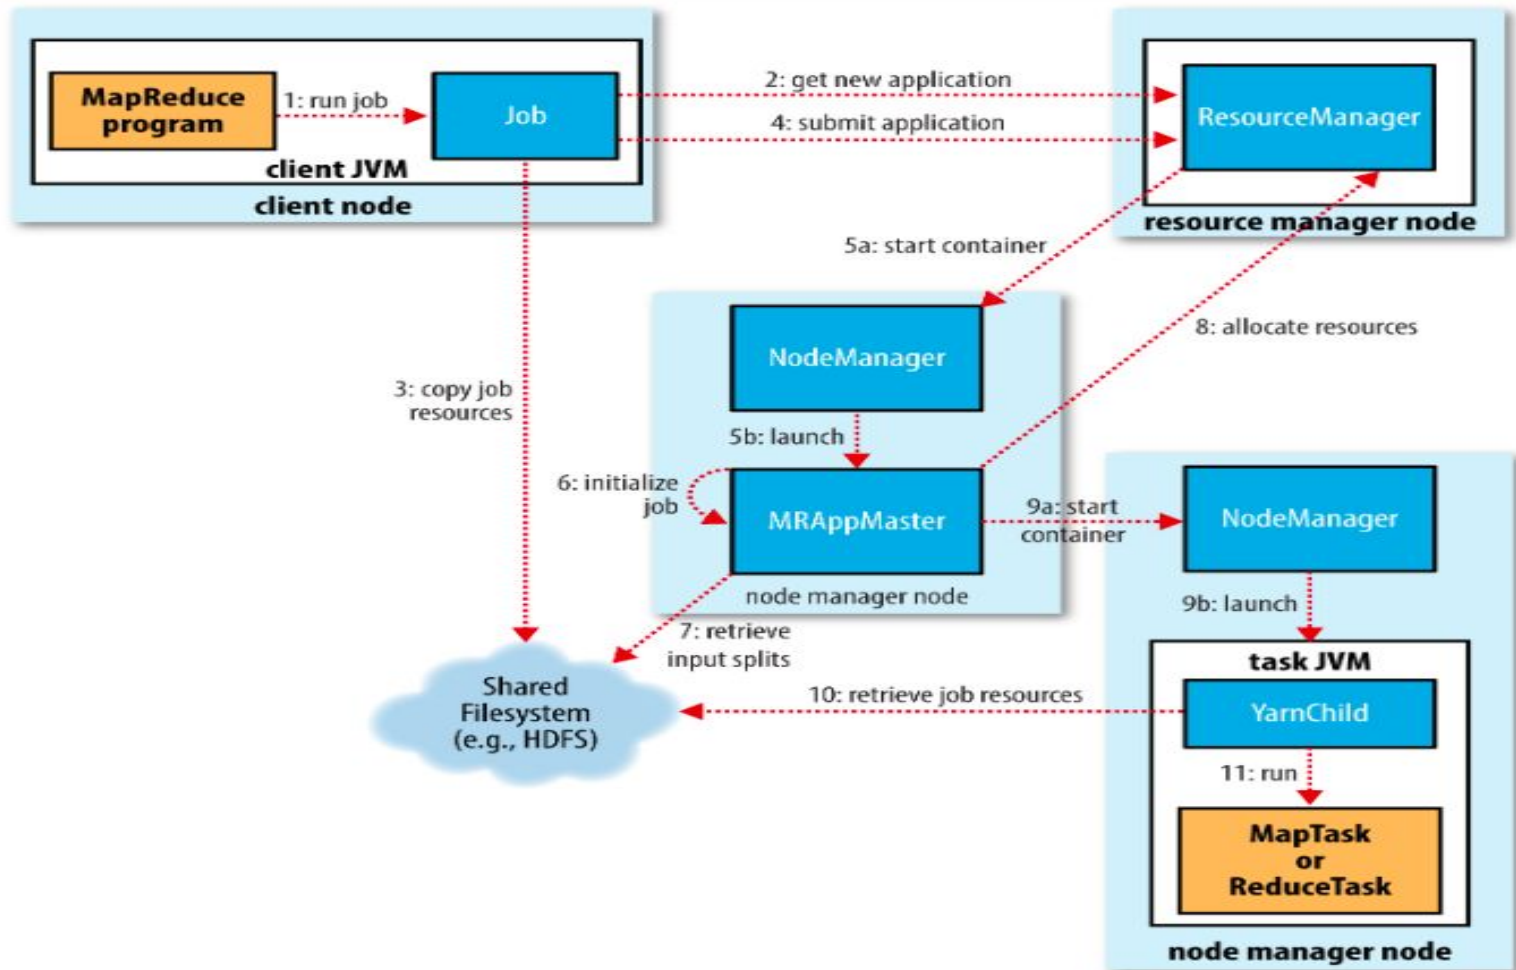


YARN Architecture

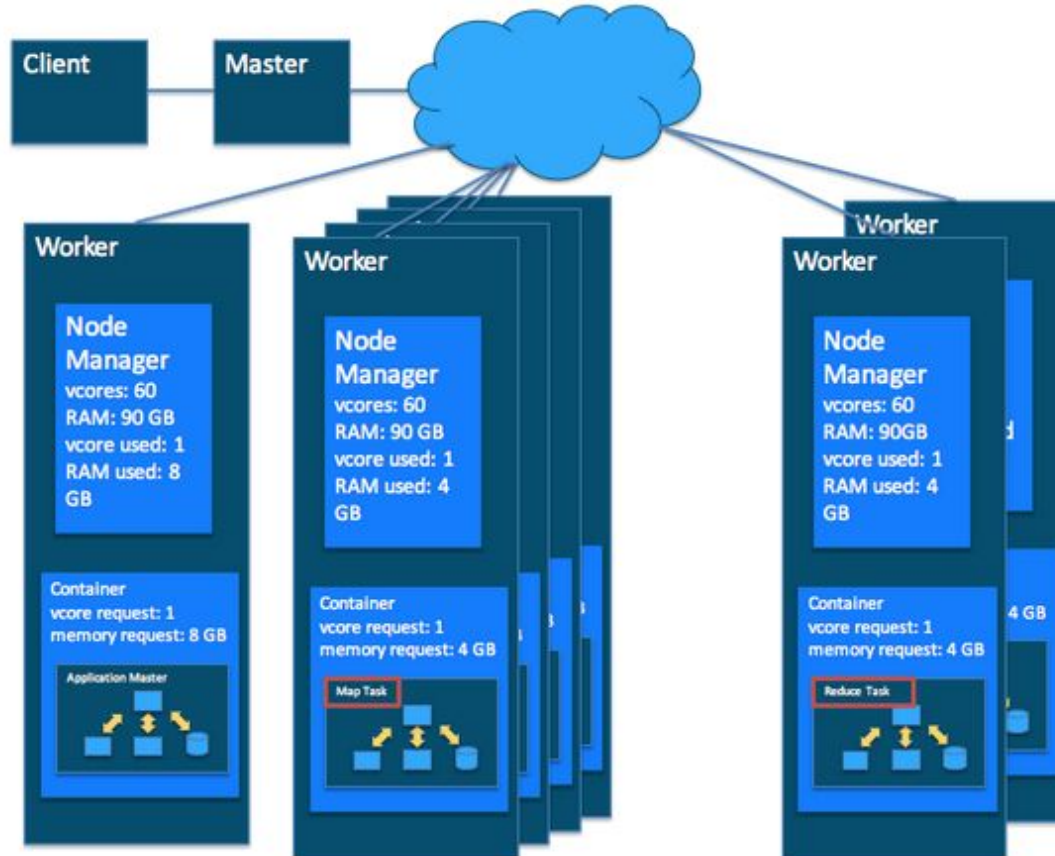


MapReduce Architecture

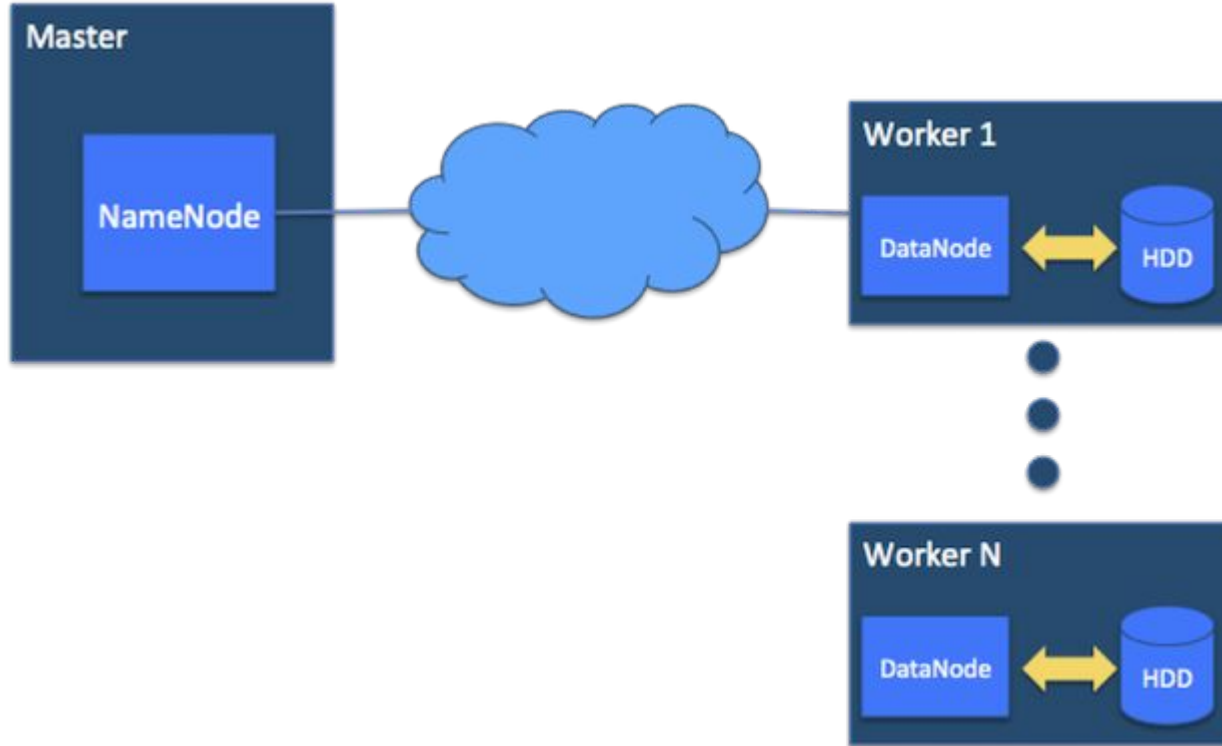





MR Application in Yarn



Resources Not Available to Yarn Apps




Fair Scheduler and Queues



```
<?xml version="1.0"?>
<allocations>
  <queue name="marketing">
    <weight>30.0</weight>
  </queue>
  <queue name="sales">
    <weight>20.0</weight>
  </queue>
  <queue name="datascience">
    <weight>40.0</weight>
  </queue>

  <queue name="admin">
    <aclSubmitApps>fred,greg</aclSubmitApps>
    <weight>10.0</weight>
  </queue>
</allocations>
```



```
<?xml version="1.0"?>
<allocations>
  <queue name="marketing">
    <weight>3.0</weight>
    <queue name="reports">
      <weight>40.0</weight>
    </queue>
    <queue name="website">
      <weight>20.0</weight>
    </queue>
  </queue>
  <queue name="sales">
    <weight>4.0</weight>
    <queue name="northamerica">
      <weight>30.0</weight>
    </queue>
    <queue name="europe">
      <weight>30.0</weight>
    </queue>
  </queue>
  <queue name="datascience">
    <weight>13.0</weight>
    <queue name="short_jobs">
      <weight>100.0</weight>
    </queue>
    <queue name="best_effort_jobs">
      <weight>0.0</weight>
    </queue>
  </queue>
</allocations>
```

```
<queue name="datascience">  
  <maxRunningApps>16</maxRunningApps>  
  <minResources>  
    6799360 mb, 1660vcores  
  </minResources>  
  <maxResources>  
    12238848 mb, 2988vcores  
  </maxResources>  
  <weight>4.0</weight>  
  <schedulingPolicy>fair</schedulingPolicy>  
</queue>
```